What

This is a package that does/has only one thing: the complete transcriptions of all episodes of The Office! (US version).

Use this data set to master NLP or text analysis. Let's scratch the surface of the subject with a few examples from the excellent Text Mining with R book, by Julia Silge and David Robinson.

First install the package from CRAN:

```
# install.packages("schrute")
library(schrute)
```

There is only one data set with the schrute package; assign it to a variable

```
mydata <- schrute::theoffice
```

Take a peek at the format:

```
dplyr::glimpse(mydata)
#> Observations: 55,130
#> Variables: 7
#> $ index             1, 358, 715, 1072, 1429, 1786, 2143, 2500, 2857…
#> $ season            "01", "01", "01", "01", "01", "01", "01", "01",…
#> $ episode           "01", "01", "01", "01", "01", "01", "01", "01",…
#> $ episode_name      " Pilot", " Pilot", " Pilot", " Pilot", " Pilot…
#> $ character         "Michael", "Jim", "Michael", "Jim", "Michael", …
#> $ text              " All right Jim. Your quarterlies look very goo…
#> $ text_w_direction  " All right Jim. Your quarterlies look very goo…
```

```
 mydata %>%
  dplyr::filter(season == '01') %>%
  dplyr::filter(episode == '01') %>%
  dplyr::slice(1:3) %>%
  knitr::kable()
```

index	season	episode	episode_name	character	text	text_w_direction
1	01	01	Pilot	Michael	All right Jim. Your quarterlies look very good. How are things at the library?	All right Jim. Your quarterlies look very good. How are things at the library?
358	01	01	Pilot	Jim	Oh, I told you. I couldn't close it. So…	Oh, I told you. I couldn't close it. So…
715	01	01	Pilot	Michael	So you've come to the master for guidance? Is this what you're saying, grasshopper?	So you've come to the master for guidance? Is this what you're saying, grasshopper?

So what we have is the season, episode number and name, character, the line spoken and the line spoken with the stage direction (cue).

We can tokenize all of the lines with a few lines from the tidytext package:

```
token.mydata <- mydata %>%
  tidytext::unnest_tokens(word, text)
```

This increases our data set to 575146 records, where each record contains a word from the script.

```
 token.mydata %>%
```

```
  dplyr::filter(season == '01') %>%
  dplyr::filter(episode == '01') %>%
  dplyr::slice(1:3) %>%
  knitr::kable()
```

index	season	episode	episode_name	character	text_w_direction	word
1	01	01	Pilot	Michael	All right Jim. Your quarterlies look very good. How are things at the library?	all
1	01	01	Pilot	Michael	All right Jim. Your quarterlies look very good. How are things at the library?	right
1	01	01	Pilot	Michael	All right Jim. Your quarterlies look very good. How are things at the library?	jim

If we want to analyze the entire data set, we need to remove some stop words first:

```
stop_words <- tidytext::stop_words

tidy.token.mydata <- token.mydata %>%
  dplyr::anti_join(stop_words, by = "word")
```

And then see what the most common words are:

```
tidy.token.mydata %>%
  dplyr::count(word, sort = TRUE)
#> # A tibble: 19,225 x 2
#>    word        n
#>
#>  1 yeah     2895
#>  2 hey      2189
#>  3 michael  2054
#>  4 dwight   1540
#>  5 uh       1433
#>  6 gonna    1365
#>  7 jim      1345
#>  8 pam      1168
#>  9 time     1129
#> 10 guys      933
#> # … with 19,215 more rows
```

```
tidy.token.mydata %>%
  dplyr::count(word, sort = TRUE) %>%
  dplyr::filter(n > 400) %>%
  dplyr::mutate(word = stats::reorder(word, n)) %>%
  ggplot2::ggplot(ggplot2::aes(word, n)) +
  ggplot2::geom_col() +
  ggplot2::xlab(NULL) +
  ggplot2::coord_flip() +
  ggplot2::theme_minimal()
```

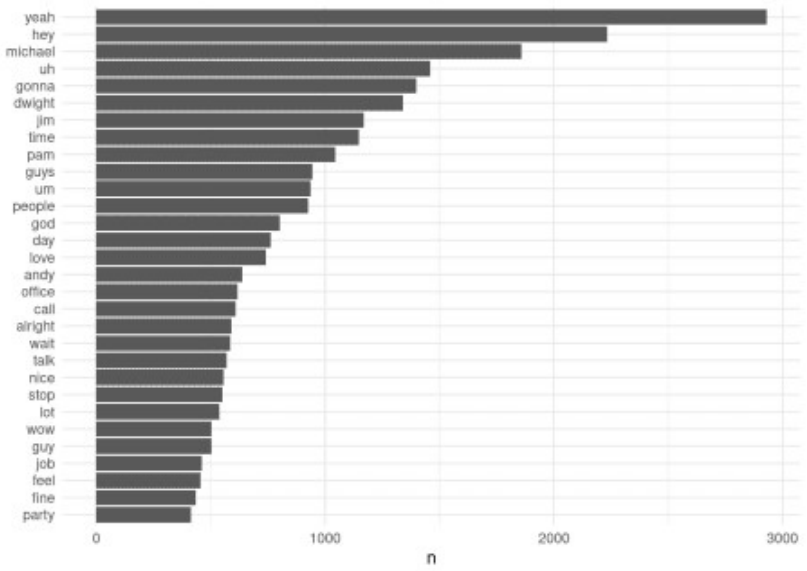

Feel free to keep going with this. Now that you have the time line (episode, season) and the character for each line and word in the series, you can perform an unlimited number of analyses. Some ideas:

– Sentiment by character
– Sentiment by character by season
– Narcissism by season (ahem.. Nard Dog season 8-9)
– Lines by character
– Etc.